

In [1]:

```
import pandas as pd
import numpy as np
```

In [2]:

```
import matplotlib.pyplot as plt
import seaborn as sns
```

In [3]:

```
%matplotlib inline
```

In [4]:

```
df=pd.read_csv("USA_Housing.csv")
```

In [5]:

```
df.head()
```

Out[5]:

| | Avg. Area Income | Avg. Area House Age | Avg. Area Number of Rooms | Avg. Area Number of Bedrooms | Area Population | Price | Addr |
|---|---------------------|------------------------------|---------------------------------------|------------------------------------|--------------------|--------------|---|
| 0 | 79545.458574 | 5.682861 | 7.009188 | 4.09 | 23086.800503 | 1.059034e+06 | 208 Michael Ferry 674\nLaurabury, 371 |
| 1 | 79248.642455 | 6.002900 | 6.730821 | 3.09 | 40173.072174 | 1.505891e+06 | 188 Johnson Vi Suite 079\nL Kathleen, C |
| 2 | 61287.067179 | 5.865890 | 8.512727 | 5.13 | 36882.159400 | 1.058988e+06 | 9127 Elizal Stravenue\nDanielc WI 0641 |
| 3 | 63345.240046 | 7.188236 | 5.586729 | 3.26 | 34310.242831 | 1.260617e+06 | USS Barnett\nFPC 44 |
| 4 | 59982.197226 | 5.040555 | 7.839388 | 4.23 | 26354.109472 | 6.309435e+05 | USNS Raymond\nF AE 09 |

In [6]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 7 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   Avg. Area Income                     5000 non-null   float64
 1   Avg. Area House Age                  5000 non-null   float64
 2   Avg. Area Number of Rooms            5000 non-null   float64
 3   Avg. Area Number of Bedrooms         5000 non-null   float64
 4   Area Population                      5000 non-null   float64
 5   Price                               5000 non-null   float64
 6   Address                             5000 non-null   object
dtypes: float64(6), object(1)
memory usage: 273.6+ KB
```

In [7]:

```
df.describe()
```

Out[7]:

| | Avg. Area Income | Avg. Area House Age | Avg. Area Number of Rooms | Avg. Area Number of Bedrooms | Area Population | Price |
|-------|------------------|---------------------|---------------------------|------------------------------|-----------------|--------------|
| count | 5000.000000 | 5000.000000 | 5000.000000 | 5000.000000 | 5000.000000 | 5.000000e+03 |
| mean | 68583.108984 | 5.977222 | 6.987792 | 3.981330 | 36163.516039 | 1.232073e+06 |
| std | 10657.991214 | 0.991456 | 1.005833 | 1.234137 | 9925.650114 | 3.531176e+05 |
| min | 17796.631190 | 2.644304 | 3.236194 | 2.000000 | 172.610686 | 1.593866e+04 |
| 25% | 61480.562388 | 5.322283 | 6.299250 | 3.140000 | 29403.928702 | 9.975771e+05 |
| 50% | 68804.286404 | 5.970429 | 7.002902 | 4.050000 | 36199.406689 | 1.232669e+06 |
| 75% | 75783.338666 | 6.650808 | 7.665871 | 4.490000 | 42861.290769 | 1.471210e+06 |
| max | 107701.748378 | 9.519088 | 10.759588 | 6.500000 | 69621.713378 | 2.469066e+06 |

In [8]:

```
df.columns
```

Out[8]:

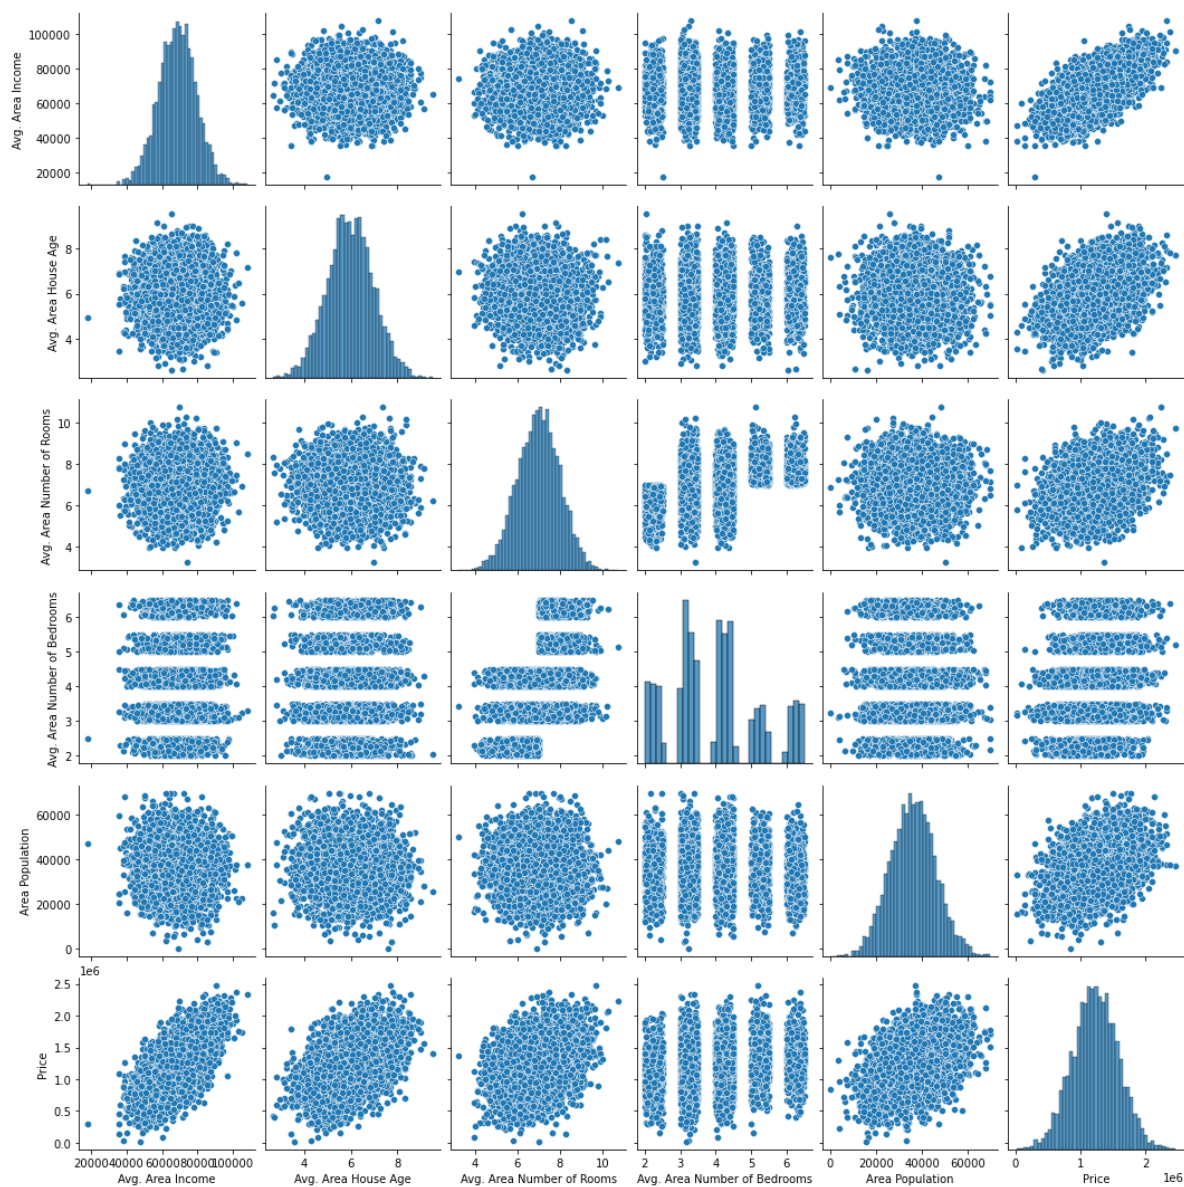
```
Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',
      'Avg. Area Number of Bedrooms', 'Area Population', 'Price', 'Address'],
      dtype='object')
```

In [9]:

```
sns.pairplot(df)
```

Out[9]:

<seaborn.axisgrid.PairGrid at 0x1d52f87a520>



In [10]:


```
sns.distplot(df["Price"])
```

C:\Users\Krishna Reddy\anaconda3\lib\site-packages\seaborn\distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[10]:

<AxesSubplot:xlabel='Price', ylabel='Density'>



In [11]:



```
df.corr()
```

Out[11]:


| | Avg. Area Income | Avg. Area House Age | Avg. Area Number of Rooms | Avg. Area Number of Bedrooms | Area Population | Price |
|------------------------------------|---------------------|------------------------|---------------------------------|------------------------------------|--------------------|----------|
| Avg. Area Income | 1.000000 | -0.002007 | -0.011032 | 0.019788 | -0.016234 | 0.639734 |
| Avg. Area House Age | -0.002007 | 1.000000 | -0.009428 | 0.006149 | -0.018743 | 0.452543 |
| Avg. Area Number of Rooms | -0.011032 | -0.009428 | 1.000000 | 0.462695 | 0.002040 | 0.335664 |
| Avg. Area Number of Bedrooms | 0.019788 | 0.006149 | 0.462695 | 1.000000 | -0.022168 | 0.171071 |
| Area Population | -0.016234 | -0.018743 | 0.002040 | -0.022168 | 1.000000 | 0.408556 |
| Price | 0.639734 | 0.452543 | 0.335664 | 0.171071 | 0.408556 | 1.000000 |

In [12]:

```
sns.heatmap(df.corr(), annot=True)
```

Out[12]:

<AxesSubplot:>



In [13]:

```
df.columns
```

Out[13]:

```
Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',  
      'Avg. Area Number of Bedrooms', 'Area Population', 'Price', 'Address',  
      ],  
      dtype='object')
```

In [14]:

```
from sklearn.model_selection import train_test_split
```

In [15]:

```
x=df[['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',  
      'Avg. Area Number of Bedrooms', 'Area Population']]
```

In [16]:

```
y=df[['Price']]
```

In [17]:



```
x_train,x_test,y_train,y_test=train_test_split(x,y ,test_size=0.4, random_state=101)
```

In [18]:



```
from sklearn.linear_model import LinearRegression
```

In [19]:



```
lm=LinearRegression()
```

In [20]:



```
lm.fit(x_train, y_train)
```

Out[20]:

```
LinearRegression()
```

In [21]:



```
print(lm.intercept_)
```

```
[-2640159.79685191]
```

In [22]:



```
x.columns
```

Out[22]:

```
Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Room  
s',  
      'Avg. Area Number of Bedrooms', 'Area Population'],  
      dtype='object')
```

In [23]:



```
print(lm.coef_)
```

```
[[2.15282755e+01 1.64883282e+05 1.22368678e+05 2.23380186e+03  
 1.51504200e+01]]
```

In [24]:




```
prediction=lm.predict(x_test)
```

In [25]:

```
plt.scatter(prediction,y_test)
```

Out[25]:

<matplotlib.collections.PathCollection at 0x1d537aabd30>



In []:

In [26]:

```
from sklearn import metrics as mt
```

In []:

```
mt.mean_absolute_error(y_test,predictions)
```

In [39]:

```
mt.mean_absolute_error(y_test,prediction)
```

Out[39]:

82288.22251914957

In [40]:



```
a=mt.mean_squared_error(y_test,prediction)
a
```

Out[40]:

10460958907.209507

In [41]:



```
np.sqrt(a)
```

Out[41]:

102278.82922291156

In [42]:



```
mt.explained_variance_score(y_test,prediction)
```

Out[42]:

0.9178179926151797

In []:



In []:

